

Variational Auto Encoder (VAE)

Taewook Nam
MLAI, KAIST

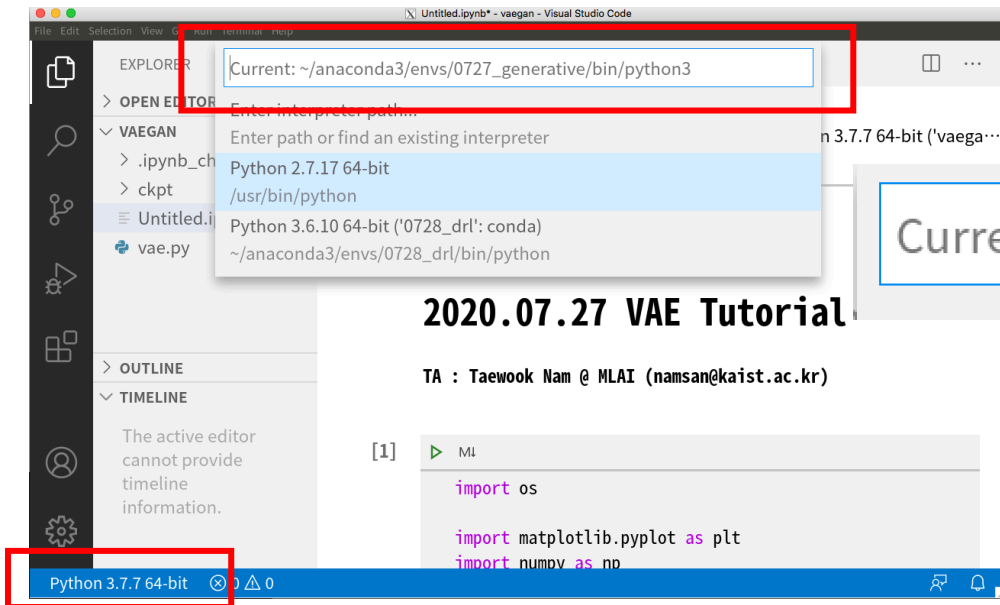
Environment

- Jupyter notebook in vscode

\$ code

- conda path

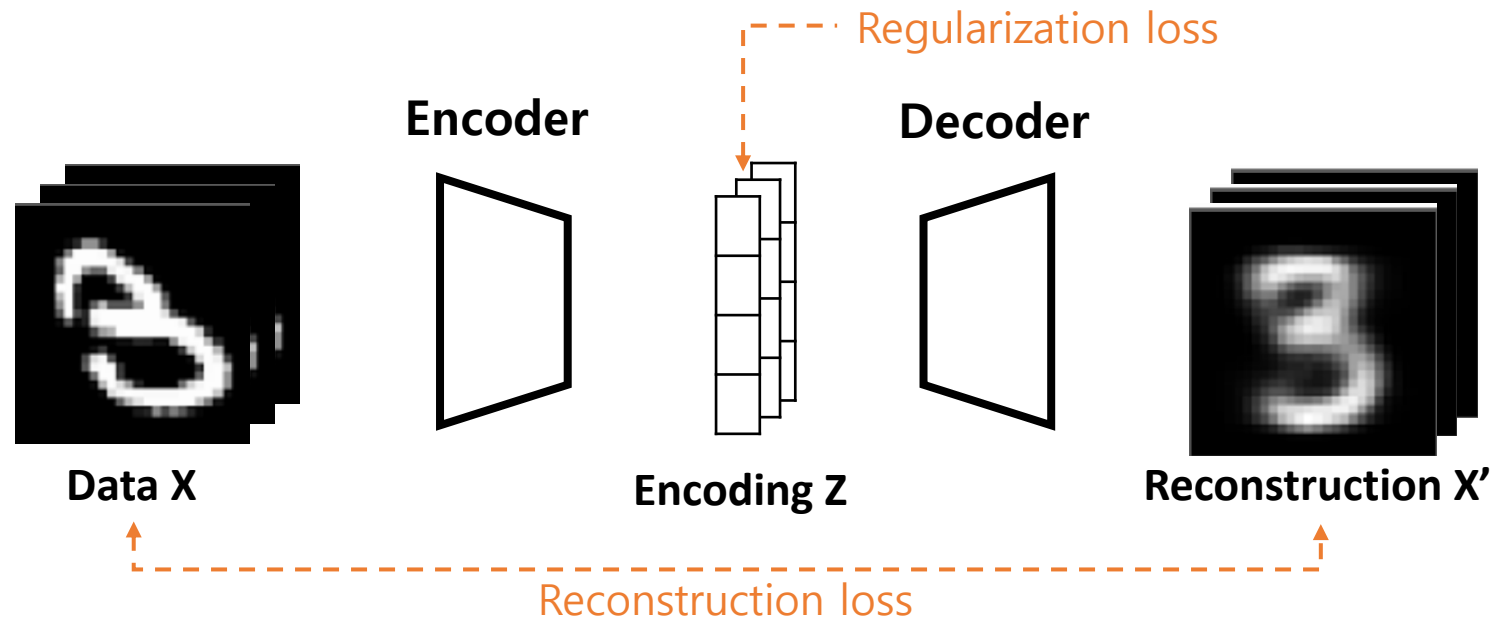
`~/anaconda3/envs/0727_generative/bin/python`



VAE Review

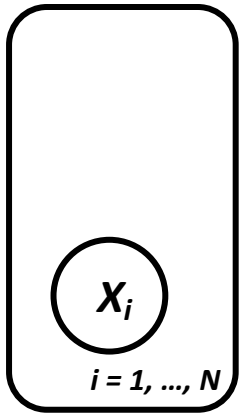
Intuitive View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*
- Why?
 - \mathbf{Z} is useful! Generator ($\mathbf{z} \rightarrow \mathbf{x}$) is useful!
- *How?*



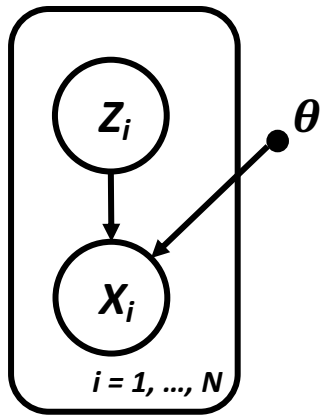
Mathematical View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$



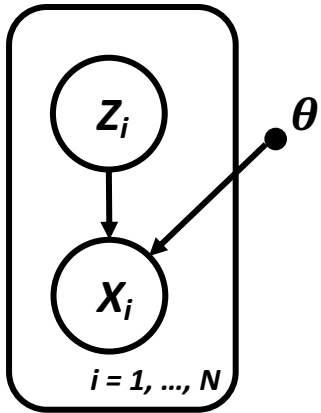
Mathematical View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*



Mathematical View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*
- **How?**

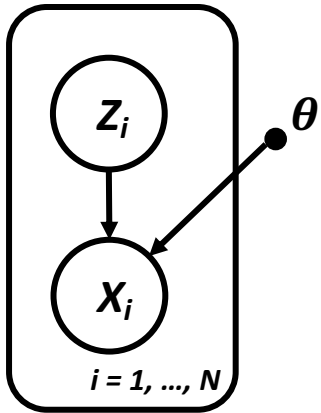


1. Directly maximizing evidence $p(\mathbf{X}|\theta)$

$$\max_{\theta} \prod_{i=1}^N p(\mathbf{x}_i|\theta) = \max_{\theta} \prod_{i=1}^N \int p(\mathbf{x}_i|\mathbf{z}, \theta) p(\mathbf{z}) d\mathbf{z}$$

Mathematical View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*
- **How?**



2. If we know true posterior?

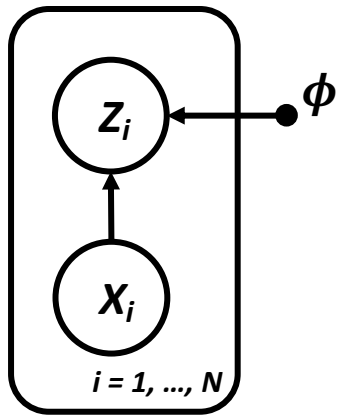
$\mathbf{z}_i \sim p(\mathbf{z}|\mathbf{x}, \theta)$ is just possible.

But this is *intractable* to compute from $p(\mathbf{x}|\mathbf{z}, \theta)$

$$p(\mathbf{z}|\mathbf{x}, \theta) = \frac{p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\theta)}{p(\mathbf{x}|\theta)} = \frac{p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\theta)}{\int p(\mathbf{x}|\mathbf{z}, \theta)p(\mathbf{z}|\theta) d\mathbf{z}}$$

Mathematical View of VAE

- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*
- **How?**

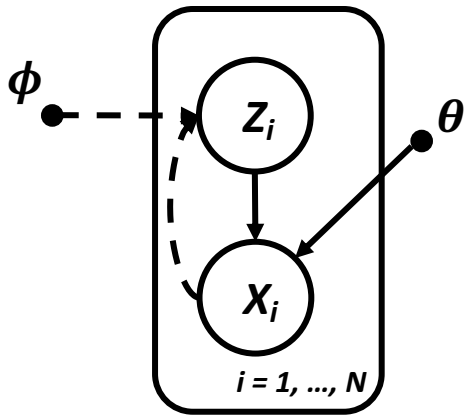


3. Model posterior?

$$p(\mathbf{x}|\phi) = \int p(\mathbf{x}|\mathbf{z}, \phi)p(\mathbf{z}|\phi) d\mathbf{z}$$

Mathematical View of VAE

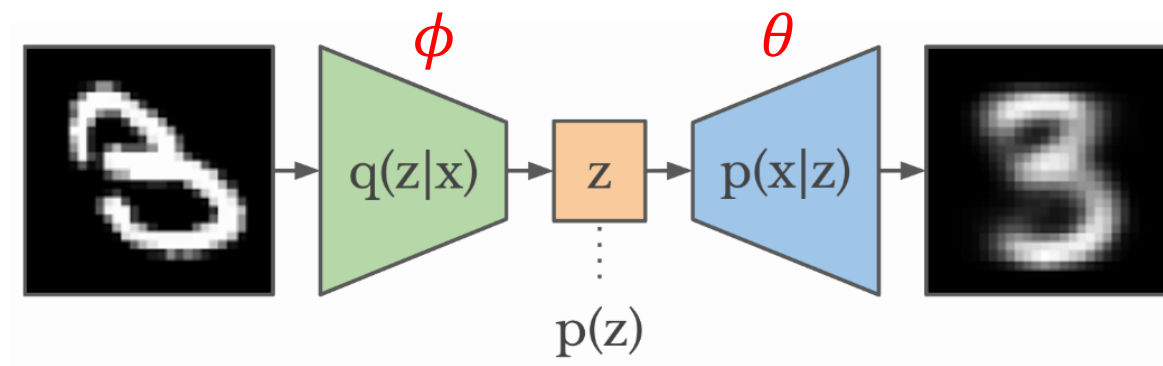
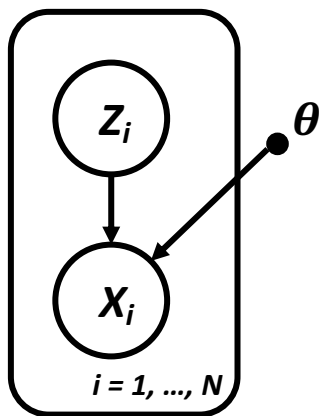
- What we have?
 - Data $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- What we want?
 - Latent variable $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$, *automatically*
- **How?**



4. “Variational Inference” does both : maximize evidence & inference posterior

$$\begin{aligned}\log p(\mathbf{x}|\theta) &= \log \int p(\mathbf{x}|\mathbf{z}, \theta) p(\mathbf{z}) d\mathbf{z} \\ &= \log \int q(\mathbf{z}|\mathbf{x}, \phi) \frac{p(\mathbf{x}|\mathbf{z}, \theta) p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x}, \phi)} d\mathbf{z} \\ &\geq \int q(\mathbf{z}|\mathbf{x}, \phi) \log \frac{p(\mathbf{x}|\mathbf{z}, \theta) p(\mathbf{z})}{q(\mathbf{z}|\mathbf{x}, \phi)} d\mathbf{z} \\ &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \phi)} [\log p(\mathbf{x}|\mathbf{z}, \theta)] - KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})]\end{aligned}$$

VAE



$$\max_{\theta, \phi} \mathbb{E}_{q(z|x, \phi)} [\log p(x|z, \theta)] - KL[q(z|x, \phi) || p(z)]$$

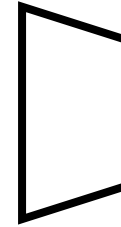
Today's Goal!

Today's Goal

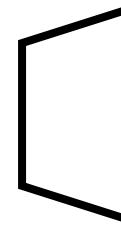
- Tensorflow implementation of VAE
- Visualization of learned encoding
- Visualization of learned decoding
- Visualization of learned latent space

```
1 from layers import *
2
3 def encoder(x, zdim, name='encoder', reuse=None):
4     x = dense(x, 500, activation=relu, name=name+'/dense1', reuse=reuse)
5     x = dense(x, 500, activation=relu, name=name+'/dense2', reuse=reuse)
6     mu = dense(x, zdim, name=name+'/mu', reuse=reuse)
7     sigma = dense(x, zdim, activation=softplus, name=name+'/sigma',
8                   reuse=reuse)
9     return mu, sigma
10
11 def decoder(x, name='decoder', reuse=None):
12     x = dense(x, 500, activation=relu, name=name+'/dense1', reuse=reuse)
13     x = dense(x, 500, activation=relu, name=name+'/dense2', reuse=reuse)
14     x = dense(x, 784, activation=sigmoid, name=name+'/output',
15              reuse=reuse)
16     return x
17
18 def autoencoder(x, zdim, training, name='autoencoder', reuse=None):
19     mu, sigma = encoder(x, zdim, reuse=reuse)
20     z = Normal(mu, sigma).sample() if training else mu
21     x_hat = decoder(z, reuse=reuse)
22
23     log_likelihood = tf.reduce_sum(x*log(x_hat) + (1-x)*log(1-x_hat), 1)
24     kl = 0.5 * tf.reduce_sum(mu**2 + sigma**2 - log(sigma**2) - 1, 1)
25     elbo = tf.reduce_mean(log_likelihood - kl)
26
27     net = {}
28     net['elbo'] = elbo
29     net['weights'] = tf.trainable_variables()
30     return net
```

Encoder

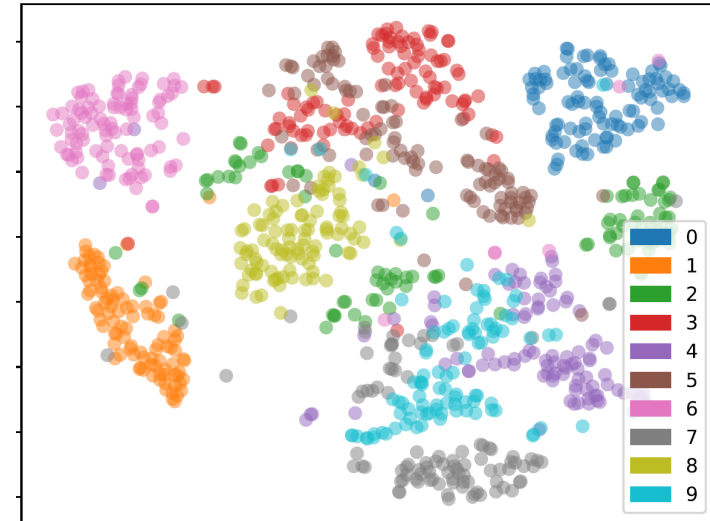
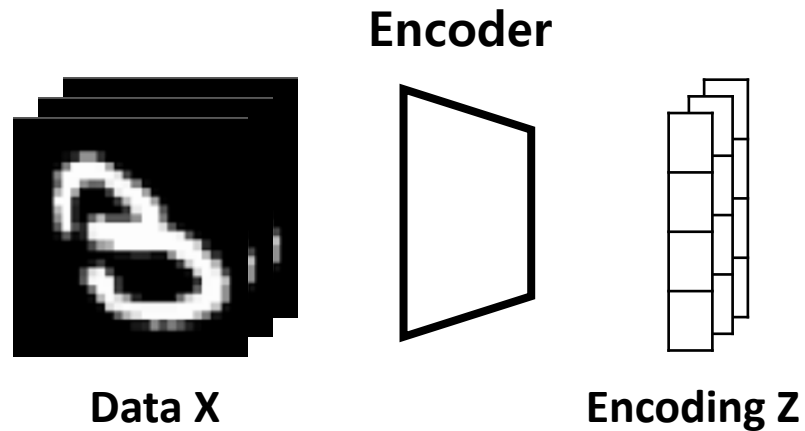


Decoder



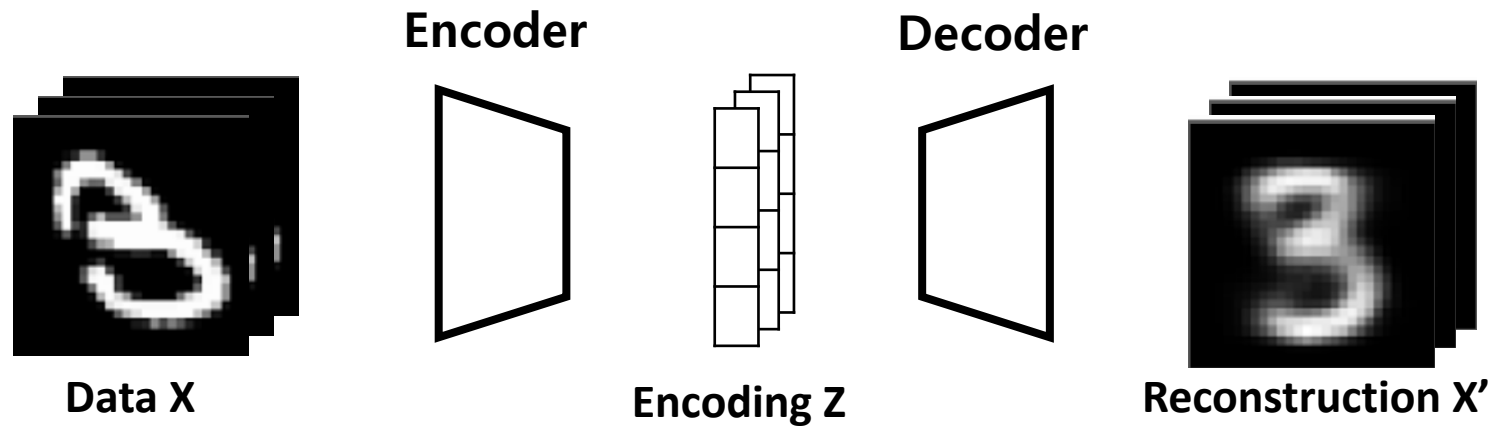
Today's Goal

- Tensorflow implementation of VAE
- Visualization of learned encoding
- Visualization of learned decoding
- Visualization of learned latent space



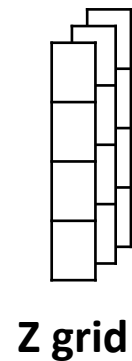
Today's Goal

- Tensorflow implementation of VAE
- Visualization of learned encoding
- Visualization of learned decoding
- Visualization of learned latent space

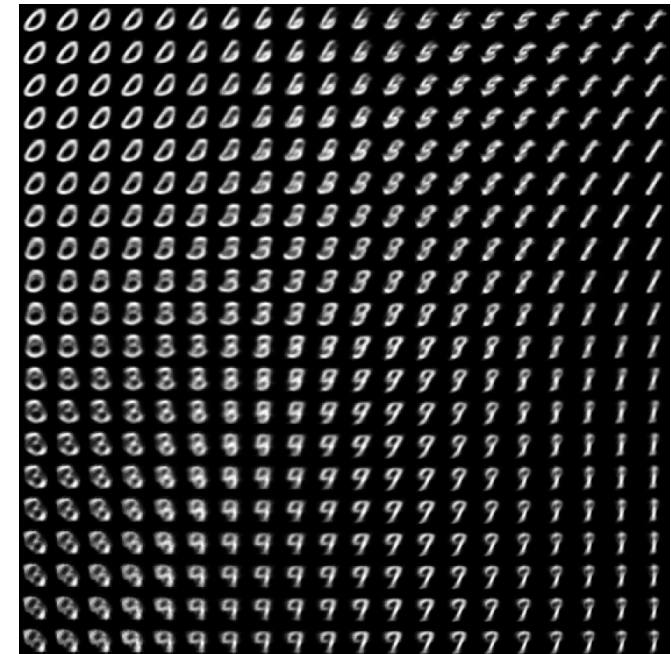


Today's Goal

- Tensorflow implementation of VAE
- Visualization of learned encoding
- Visualization of learned decoding
- Visualization of learned latent space



Decoder



Today's Goal!

Task 1 : Tensorflow implementation of VAE

Task 2 : Visualization of learned encoding

Task 3 : Visualization of learned decoding

Task 4 : Visualization of learned latent space

Task 1 : TF implementation

- Network

```
class VAE(Model):
    h_dim = 500

    def __init__(self, x_shape, z_dim):
        super().__init__()

        x_dim = np.prod(x_shape)
        self.z_dim = z_dim

        self.encoder = Sequential([
            Flatten(),
            Dense(self.h_dim, activation='relu'),
            Dense(self.h_dim, activation='relu'),
            Dense(2*z_dim, activation='relu')
        ])
        self.decoder = Sequential([
            Dense(self.h_dim, activation='relu'),
            Dense(self.h_dim, activation='relu'),
            Dense(x_dim, activation='sigmoid'),
            Reshape(x_shape)
        ])
    )
```

```
def encode(self, x):
    z_param = self.encoder(x)
    z_mu, z_rho = z_param[:, :self.z_dim], z_param[:, self.z_dim:]
    return z_mu, tf.nn.softplus(z_rho)
```

```
def decode(self, z):
    return self.decoder(z)
```

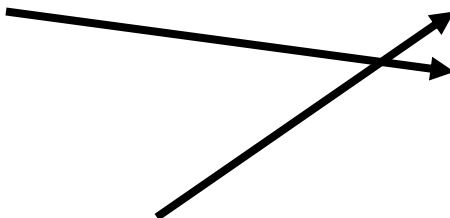
Task 1 : TF implementation

$$\max_{\theta, \phi} \mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \phi)} [\log p(\mathbf{x}|\mathbf{z}, \theta)] - KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})]$$

- MC Sampling

$$\mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \phi)} [\log p(\mathbf{x}|\mathbf{z}, \theta)] \approx \frac{1}{S} \sum_{s=1}^S \log p(\mathbf{x}|\mathbf{z}^{(s)}, \theta)$$

where $\mathbf{z}^{(s)} \sim q(\mathbf{z}|\mathbf{x}, \phi)$



```
z_mu, z_sigma = vae.encode(x)
z_dist = tfp.distributions.Normal(z_mu, z_sigma)
z = z_dist.sample()
x_reconst = vae.decode(z)
```

- Reparameterization

$$p(\mathbf{x}|\mathbf{z}^{(s)}, \theta), \quad \mathbf{z}^{(s)} \sim q(\mathbf{z}|\mathbf{x}, \phi)$$

$$\rightarrow p(\mathbf{x}|\boldsymbol{\mu}_z + \boldsymbol{\epsilon} \otimes \boldsymbol{\sigma}_z), \quad \boldsymbol{\epsilon} \sim N(\mathbf{0}, I)$$

Task 1 : TF implementation

$$\max_{\theta, \phi} \mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \phi)} [\log p(\mathbf{x}|\mathbf{z}, \theta)] - KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})]$$

- Likelihood

$$\begin{aligned} \log p(\mathbf{x}|\mathbf{z}, \theta) &= \log \prod_{d=1}^{D_x} \text{Ber}(x_d|x'_d) \\ &= \sum_{d=1}^D \log(x'_d)^{x_d} (1 - x'_d)^{1-x_d} \longrightarrow \\ &= \sum_{d=1}^D x_d \log x'_d + (1 - x_d) \log(1 - x'_d) \end{aligned}$$

```
def compute_elbo(x, x_reconst, z_mu, z_sigma):  
    log_likelihood = tf.reduce_sum(  
        x * tf.math.log(x_reconst + 1e-6) + (1 - x) * tf.math.log(1 - x_reconst + 1e-6),  
        axis=(1, 2)  
    )  
    kl = 0.5 * tf.reduce_sum(  
        z_mu**2 + z_sigma**2 - tf.math.log(z_sigma**2) - 1,  
        axis=1  
    )  
    elbo = tf.reduce_mean(log_likelihood - kl)  
    return elbo
```

- KL divergence

$$\begin{aligned} KL[q(\mathbf{z}|\mathbf{x}, \phi) || p(\mathbf{z})] &= \sum_{d=1}^{D_z} KL[N(z_d|\mu_d, \sigma_d^2) || N(z_d|0, 1)] \\ &= \sum_{d=1}^{D_z} \mu_d^2 + \sigma_d^2 - \log \sigma_d^2 - 1 \end{aligned}$$

Task 1 : TF implementation

- Training

```
for epoch_i in trange(100):
    for x in train_ds:
        with tf.GradientTape() as tape:
            z_mu, z_sigma = vae.encode(x)
            z_dist = tfp.distributions.Normal(z_mu, z_sigma)
            z = z_dist.sample()
            x_reconst = vae.decode(z)

            elbo = compute_elbo(x, x_reconst, z_mu, z_sigma)
            loss = -elbo
            gradients = tape.gradient(loss, vae.trainable_variables)
            optimizer.apply_gradients(zip(gradients, vae.trainable_variables))

        if (epoch_i+1) % 5 == 0:
            vae.save_weights(f'ckpt/{epoch_i + 1}')
```

- Load

```
ckpt_i = 100
vae.load_weights(f'ckpt/{ckpt_i}')
```

Today's Goal!

Task 1 : Tensorflow implementation of VAE

Task 2 : Visualization of learned encoding

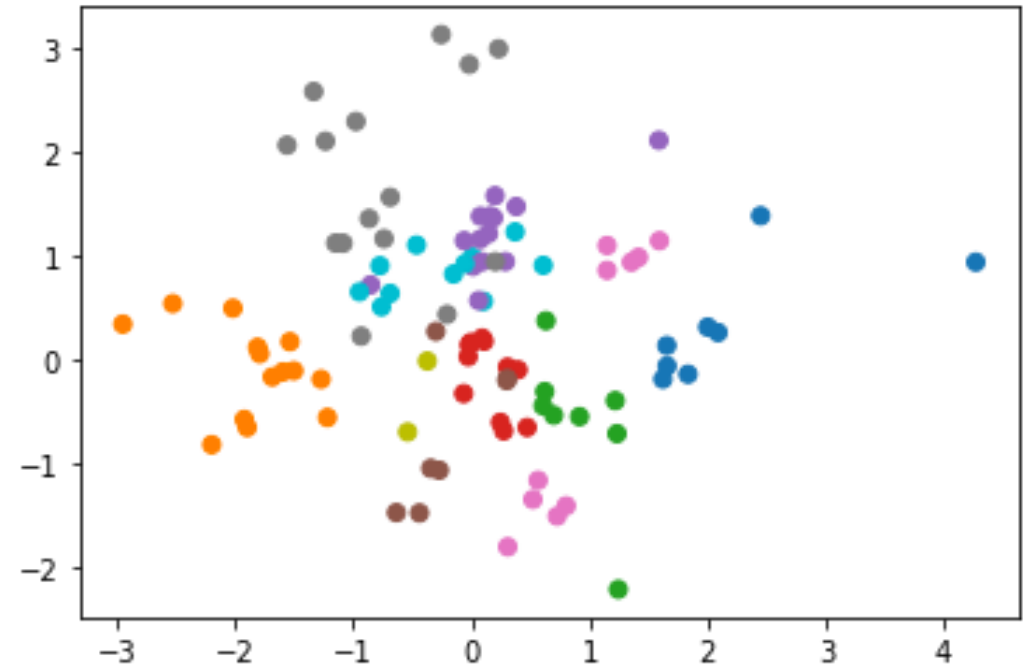
Task 3 : Visualization of learned decoding

Task 4 : Visualization of learned latent space

Task 2 : Learned Encoding

```
n_points = 100
z_mu, _ = vae.encode(x_test[:n_points])

fig, ax = plt.subplots()
scatter = plt.scatter(
    z_mu[:, 0], z_mu[:, 1],
    c=[f'C{y_test[i]}' for i in range(n_points)],
    label=[y_test[i] for i in range(n_points)]
)
```



Today's Goal!

Task 1 : Tensorflow implementation of VAE

Task 2 : Visualization of learned encoding

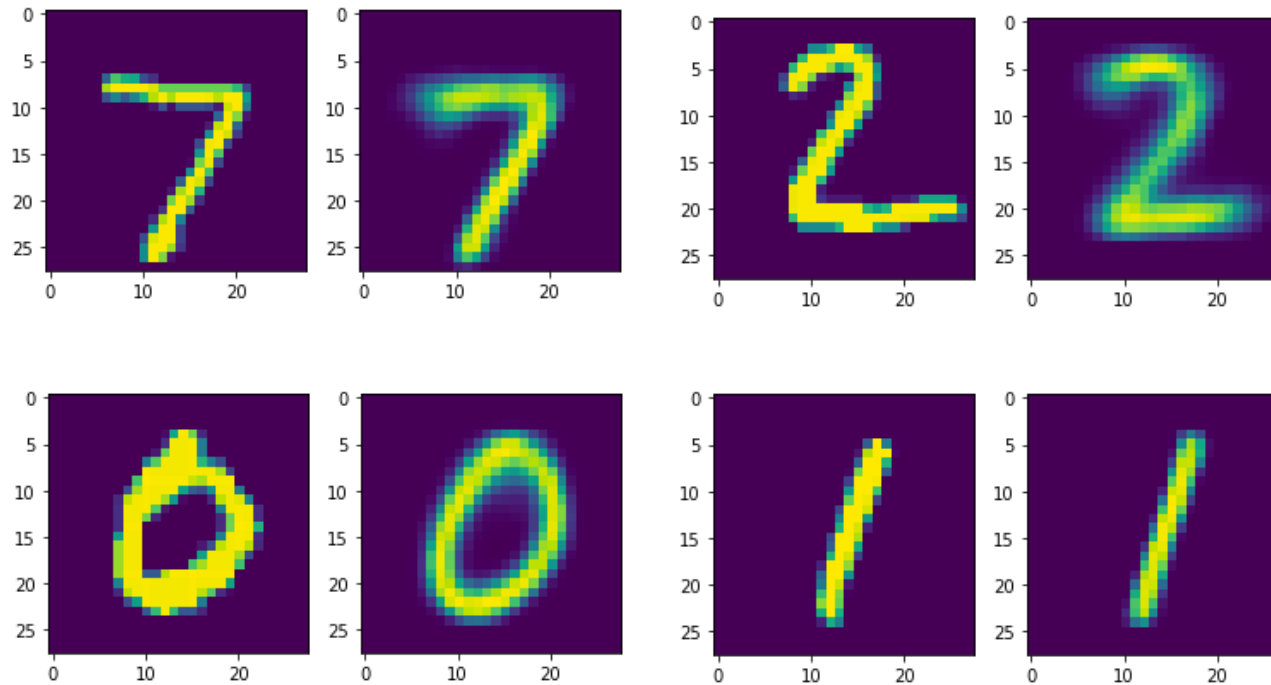
Task 3 : Visualization of learned decoding

Task 4 : Visualization of learned latent space

Task 3 : Learned Decoding

```
idx = 10
z_mu, _ = vae.encode(x_test)
x_reconst = vae.decode(z_mu)

fig, (ax1, ax2) = plt.subplots(1, 2)
ax1.imshow(x_test[idx])
ax2.imshow(x_reconst[idx])
```



Today's Goal!

Task 1 : Tensorflow implementation of VAE

Task 2 : Visualization of learned encoding

Task 3 : Visualization of learned decoding

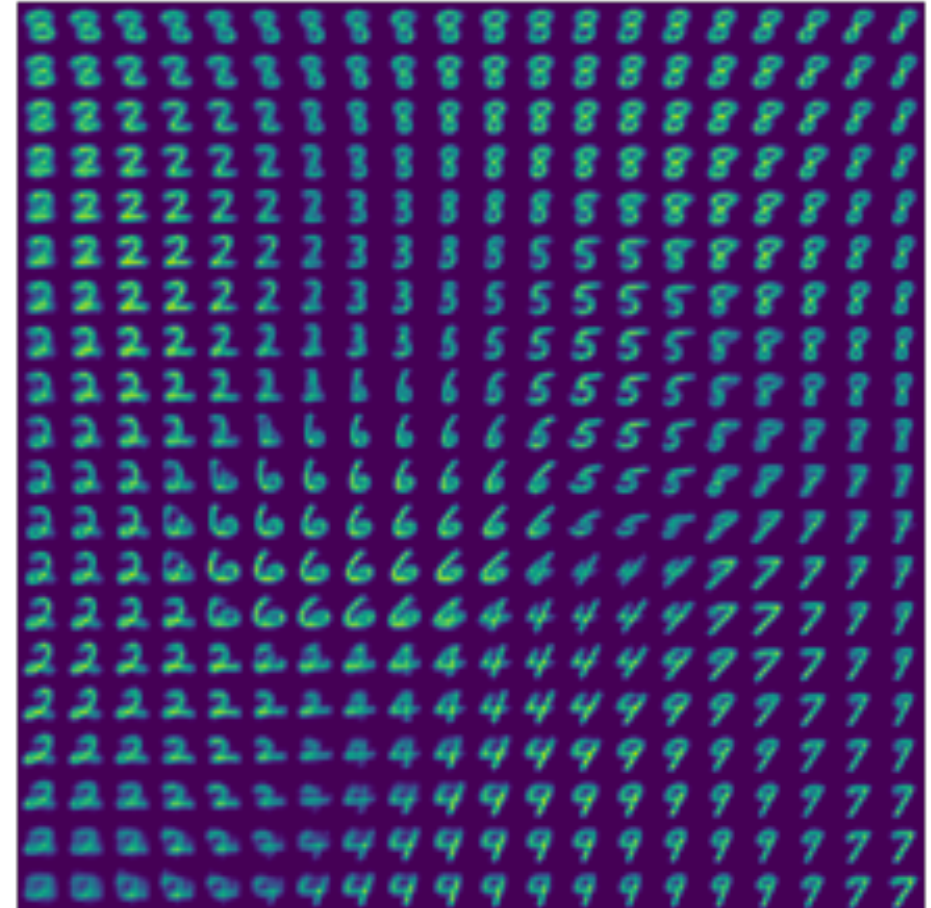
Task 4 : Visualization of learned latent space

Task 4 : Learned Latent Space

```
z_range = np.arange(-1, 1, 0.1)
z_grid = np.stack(np.meshgrid(z_range, z_range), axis=-1)

z_grid_flat = z_grid.reshape(400, 2)
x_reconst_flat = vae.decode(tf.convert_to_tensor(z_grid_flat))
x_reconst = x_reconst_flat.numpy().reshape(20, 20, 28, 28)
```

```
x_merged = x_reconst.swapaxes(1, 2).reshape(20*28, 20*28)
plt.imshow(x_merged)
plt.axis('off')
```



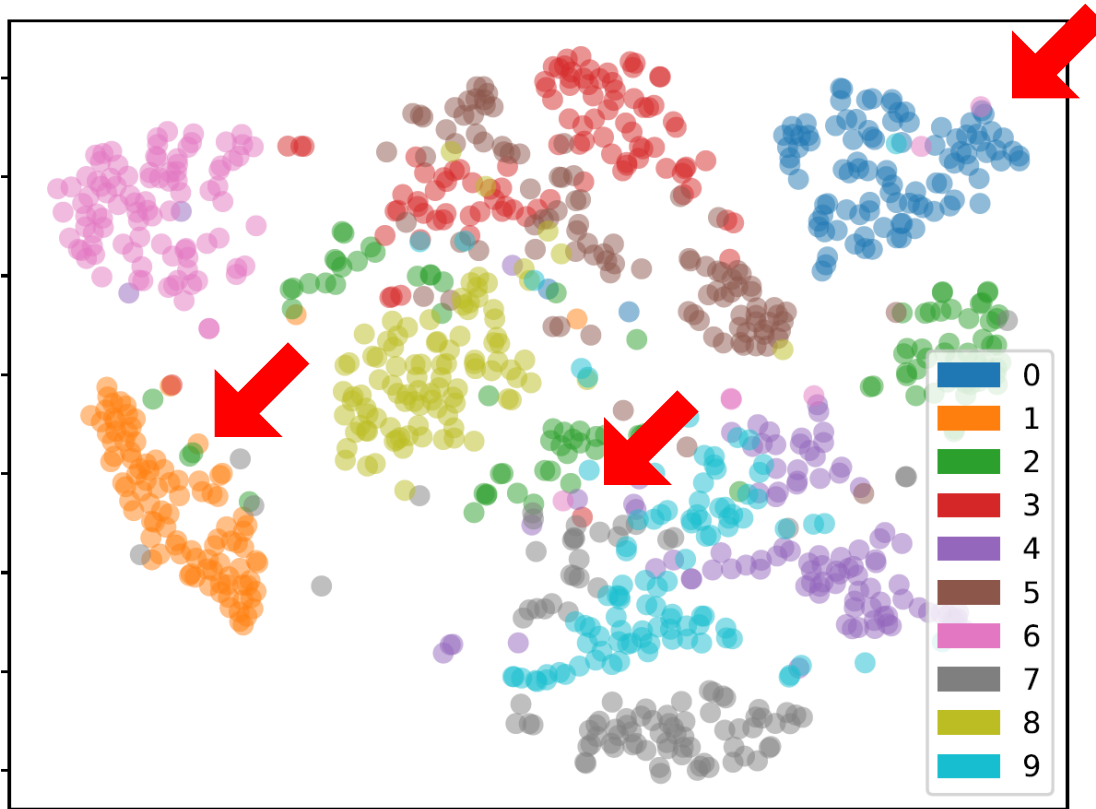
Task 4 : Learned Latent Space

- Plot

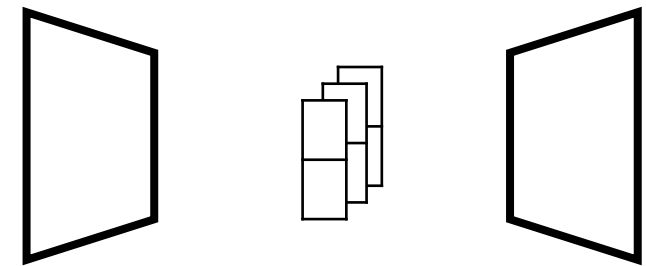
Research Questions

Research Questions

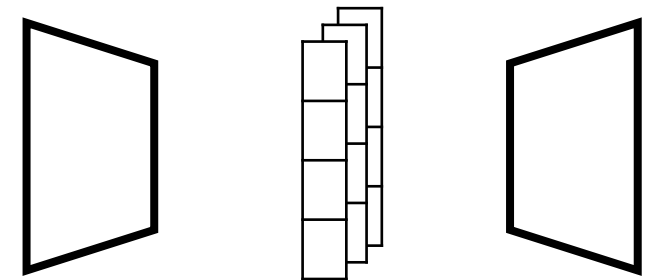
Who are they?



Difference?



Low-dimensional z



High-dimensional z

Research Questions

What happen if we omit this?

And your own 😊!!

$$\log p(X; \theta)$$

$$\geq \sum_{x \in X} E_{q(z|x; \phi)} [\log p(x|z; \theta)]$$

$$- KL[q(z|x; \phi) || p(z)]$$

감사합니다